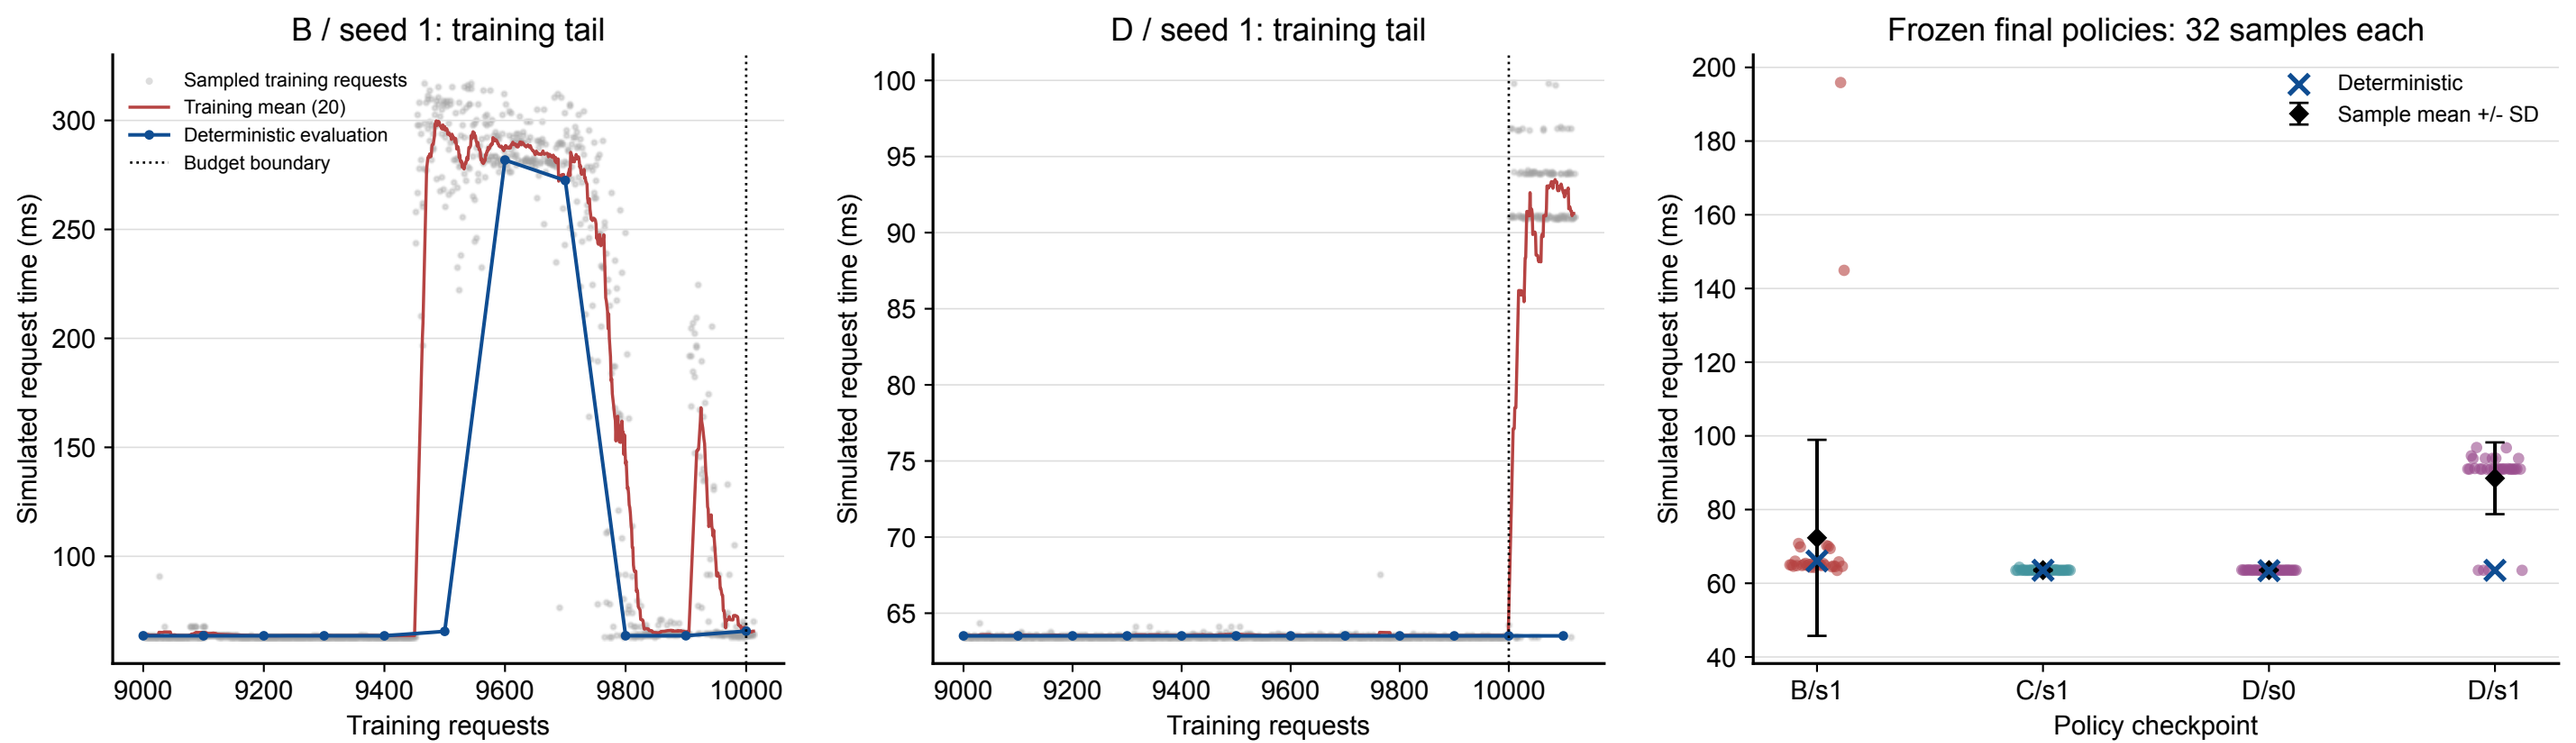
B / seed 1: training tail
Sampled training requests
Training mean (20)
Deterministic evaluation
Budget boundary
Simulated request time (ms)
Training requests
D / seed 1: training tail
Frozen final policies: 32 samples each
Deterministic
Sample mean +/- SD
B/s1
C/s1
D/s0
D/s1
Policy checkpoint